PatentForge

AI-Powered Full-Lifecycle Patent Platform

Version 0.8.3 | April 2026 | Scott Converse

PatentForge is a research tool, not a legal service. The author of this tool is not a lawyer. This tool does not provide legal advice. It helps inventors explore the patent landscape for their ideas using AI. The output is for your personal educational use and is not a substitute for professional legal counsel.

Table of Contents

1. Overview
2. What PatentForge Does
3. Important Disclaimer
4. Download and Installation
5. Quick Start (from Source)
6. System Architecture
7. The 6-Stage Feasibility Pipeline
8. AI-Assisted Claim Drafting
9. Compliance Checking
10. Patent Application Generator
11. Docker Deployment
12. Configuration
13. Legal Guardrails
14. What's New in v0.8.3
15. What's New in v0.8.2
16. Roadmap
17. License

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

1. Overview

PatentForge is a self-hosted web application that helps inventors organize their thinking about an invention before consulting a patent attorney. It uses Claude AI to structure a technical analysis, search for related patents, and highlight the questions a patent professional would explore --- so you walk into that first meeting prepared, not cold.

PatentForge is an open-source, full-lifecycle patent platform that takes an inventor from "I have an idea" to "here's a draft patent application with prior art citations, compliance checks, and a filing strategy." It is the only open-source tool that covers the entire patent lifecycle in one platform.

The platform consists of **6 services** working together:

- **React Frontend** --- TypeScript, Tailwind CSS, Vite (port 8080)
- **NestJS Backend** --- Prisma ORM, SQLite/PostgreSQL (port 3000)
- **Feasibility Service** --- TypeScript, Anthropic SSE streaming (port 3001)
- **Claim Drafter** --- Python, FastAPI, LangGraph, 3-agent pipeline (port 3002)
- **Application Generator** --- Python, FastAPI, LangGraph, 5-agent pipeline (port 3003)
- **Compliance Checker** --- Python, FastAPI, LangGraph, 4 checker agents (port 3004)

Additionally, a **System Tray App** (Go binary) manages all services in the installer builds, with health monitoring, auto-restart, and log rotation.

Test Coverage: 595 automated tests (237 backend + 142 frontend + 68 claim-drafter + 50 compliance + 64 application-generator + 34 E2E).

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

2. What PatentForge Does

- **Structured technical analysis** --- 6-stage AI pipeline that restates your invention in patent terminology, searches for related prior art, identifies potential issues under patent law, and organizes findings into a readable report
- **Prior art discovery** --- automated patent search via USPTO Open Data Portal with improved relevance scoring (stop-word filtering, title weighting), plus AI-powered web search for related patents, papers, and products
- **Patent claims viewer** --- lazy-loads actual patent claims text from the USPTO Documents API when you click a prior art result (requires free ODP API key)
- **Cost transparency** --- pre-run cost estimate based on historical run data (within 25% of actual), per-stage token tracking, and configurable cost cap
- **Resume from interruption** --- pick up where you left off if a run stops mid-pipeline
- **Multiple export formats** --- HTML, Word (.docx), and Markdown for reports; Word export for claims and compliance results; CSV for prior art
- **AI-assisted claim drafting** --- 3-agent pipeline (Planner, Writer, Examiner) generates patent claim drafts informed by your feasibility analysis and prior art, with per-claim examiner review
- **Claim tree visualization** --- SVG-based hierarchical view of claim dependencies with list/tree toggle
- **Patent family lookup** --- continuity data (parents, children, continuations, divisionals) fetched from USPTO ODP and displayed in the patent detail drawer
- **Compliance checking** --- automated validation of claim drafts against 35 USC 112(a), 35 USC 112(b), MPEP 608, and 35 USC 101 (Alice/Mayo framework), with traffic-light PASS/FAIL/WARN results, MPEP citations, and actionable fix suggestions
- **Patent application generator** --- 5-agent LangGraph pipeline assembles a complete USPTO-formatted patent application (background, summary, detailed description, abstract, figure descriptions, IDS) from feasibility analysis, prior art, and claim drafts; exports as USPTO-compliant Word (.docx) or Markdown
- **ODP API usage tracking** --- queries, results, and rate limit events tracked per search with weekly summary in Settings
- **Self-hosted** --- runs on your machine; invention data stays local except for Anthropic API calls
- **Configurable** --- choose your model (Sonnet, Opus, Haiku), set max tokens, adjust inter-stage delays

---

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

3. Important Disclaimer

PatentForge is a **research and preparation tool**, not a legal service. It does not provide legal advice, patent opinions, or attorney services. The output is intended to help you prepare for a consultation with a registered patent attorney or patent agent --- not to replace one. Patent law is complex, and decisions about whether to file a patent application should always be made with qualified legal counsel.

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

4. Download and Installation

The easiest way to get started is to download the installer for your platform:

Platform	Download	Notes
Windows	PatentForge-0.8.3-Setup.exe	Inno Setup installer (~100 MB)
Mac (Beta)	PatentForge-0.8.3.dmg	Drag to Applications (~100 MB)
Linux (Beta)	PatentForge-0.8.3.AppImage	chmod +x and run (~120 MB)

Download from: https://github.com/scottconverse/patentforge/releases/latest

The installer bundles everything --- no Node.js, Python, or git required. On first launch, the system tray icon appears and your browser opens to the setup wizard, which walks you through entering your Anthropic API key.

What the installer includes:

- Node SEA binaries for backend and feasibility service (no Node.js runtime needed)
- Portable Python 3.12 for claim-drafter, application-generator, and compliance-checker
- System tray app (Go binary) for service management
- React frontend served by backend in production mode

Mac and Linux installers are beta --- please report issues on GitHub Issues.

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

5. Quick Start (from Source)

Prerequisites

- Node.js 18+ (recommended: 20 LTS)
- Python 3.11+ (for claim-drafter, compliance-checker, and application-generator services)
- Anthropic API key (set in the Settings UI after first launch)

Install and Run

```
git clone https://github.com/scottconverse/patentforge.git
cd patentforge

# Install Node dependencies
cd backend && npm install && cd ..
cd services/feasibility && npm install && cd ..
cd frontend && npm install && cd ..

# Install Python dependencies
pip install uvicorn fastapi anthropic

# Create backend .env (SQLite, zero config)
echo 'DATABASE_URL="file:./prisma/dev.db"' > backend/.env

# Set up the database
cd backend && npx prisma db push && npx prisma generate && cd ..
```

On Windows --- double-click PatentForge.bat (builds and starts all 5 services, opens browser).

Manual start (any OS) --- run each in a separate terminal:

Terminal	Service	Command	Port
1	Backend	cd backend && npm run build && npm start	3000
2	Feasibility	cd services/feasibility && npm run build && npm start	3001
3	Claim Drafter	cd services/claim-drafter && py -m uvicorn src.server:app --port 3002	3002
4	App Generator	cd services/application-generator && py -m uvicorn src.server:app --port 3003	3003
5	Compliance Checker	cd services/compliance-checker && py -m uvicorn src.server:app --port 3004	3004
6	Frontend	cd frontend && npm run dev	8080

Open http://localhost:8080, go to Settings (gear icon), enter your Anthropic API key, and create your first project.

Troubleshooting

- **Port already in use** --- kill the process occupying the port, or change the port in the service config.

- **"No API key configured"** --- open Settings (gear icon) and enter your Anthropic API key. The key is encrypted at rest.
- **Python services won't start** --- ensure uvicorn and fastapi are installed: pip install uvicorn fastapi anthropic
- **Prisma errors** --- run cd backend && npx prisma db push && npx prisma generate to reset the database schema.

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

6. System Architecture

PatentForge uses a **federated architecture**: a central backend orchestrates independent specialized services that each own one capability. Services communicate over HTTP/JSON and can be developed, deployed, scaled, and replaced independently.

Architecture Diagram (v0.8.3)

```
+-------------------+     +-------------------+     +-------------------+
|  React Frontend   |---->|  NestJS Backend   |---->|  Feasibility      |
|  (Vite, TW CSS)   |     |  (Prisma, SQLite) |     |  Service (TS)     |
|  port 8080        |<----|  port 3000        |<----|  port 3001        |
+-------------------+ SSE +-------------------+ SSE +-------------------+
                              |                       |
                              |                       v
                              |                Anthropic Claude API
                              v
                       +-------------------+     +-------------------+
                       |  Claim Drafter    |     |  Compliance       |
                       |  (Python/         |     |  Checker          |
                       |   LangGraph)      |     |  (Python/         |
                       |  port 3002        |     |   LangGraph)      |
                       +-------------------+     |  port 3004        |
                                                 +-------------------+
                       +-------------------+
                       |  Application      |
                       |  Generator        |
                       |  (Python/         |
                       |   LangGraph)      |
                       |  port 3003        |
                       +-------------------+
```

Service Details

System Tray (Installer only) --- Go binary that launches and monitors all services, with health checks, auto-restart on crash, and log rotation.

Frontend --- React 18, TypeScript, Tailwind CSS, Vite. Served by backend in production. Communicates with backend over HTTP and Server-Sent Events (SSE).

Backend --- NestJS, Prisma ORM, SQLite (dev) / PostgreSQL (Docker). Handles project CRUD, settings management, pipeline orchestration, export engine, and proxies requests to all downstream services. Node SEA binary in installer. Security headers via Helmet. API key rate limiting via @nestjs/throttler.

Feasibility Service --- Express server, Anthropic SSE streaming, 6 prompt templates. Runs the 6-stage sequential analysis pipeline. Node SEA binary in installer.

Claim Drafter --- Python, FastAPI, LangGraph. 3-agent pipeline: Planner (claim strategy), Writer (draft claims), Examiner (per-claim review). Port 3002.

Compliance Checker --- Python, FastAPI, LangGraph. 4 specialized checker agents: 35 USC 112(a) written description, 35 USC 112(b) definiteness, MPEP 608 formalities, 35 USC 101 eligibility (Alice/Mayo). Port 3004.

Application Generator --- Python, FastAPI, LangGraph. 5-agent sequential pipeline: background, summary, detailed description, abstract, figure descriptions. Assembles complete USPTO-formatted patent application. Port 3003.

External APIs

- **Anthropic Claude API** --- powers all AI analysis, claim drafting, compliance checking, and application generation
- **USPTO Open Data Portal** --- patent search, claims text, patent family/continuity data
- **Anthropic Web Search** --- used in feasibility stages 2, 3, and 4 for grounded prior art research

Database

SQLite for local development (zero config). PostgreSQL for Docker deployments. Prisma ORM manages schema migrations. API keys encrypted at rest with AES-256-GCM.

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

7. The 6-Stage Feasibility Pipeline

PatentForge runs a 6-stage sequential analysis pipeline using the Anthropic Claude API. Each stage builds on the output of all previous stages. Stages 2, 3, and 4 use Anthropic's web search tool for grounded research.

Stage	Name	What It Does
1	Technical Intake and Restatement	Restates your invention in precise technical language
2	Prior Art Research	Searches for existing patents, papers, and products
3	Patentability Review	Identifies potential issues under 35 USC 101, 102, 103, and 112
4	Deep Dive Analysis	Domain-specific landscape analysis tailored to the invention's technical areas
5	Strategy Notes	Summarizes filing considerations, cost factors, and open questions
6	Consolidated Report	Assembles all findings into a single structured document

Note: The output of this pipeline is structured research, not a legal opinion. It is designed to help you and your patent attorney have a more productive first conversation.

Pipeline Features

- **Cost estimate before run** --- based on historical data, accurate within 25% of actual cost
- **Per-stage token tracking** --- see exactly how many tokens each stage consumed
- **Configurable cost cap** --- enforced server-side; blocks new runs when cumulative project cost exceeds the cap; mid-pipeline cancellation when a stage pushes cost over
- **Resume from interruption** --- if the pipeline stops mid-run, resume from the last completed stage
- **Individual stage re-run** --- re-run any single stage without restarting the full pipeline
- **Real-time streaming** --- see AI output token-by-token as each stage executes
- **Multiple export formats** --- HTML, Word (.docx), Markdown, CSV (for prior art)

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

8. AI-Assisted Claim Drafting

The claim drafting service uses a 3-agent LangGraph pipeline to generate patent claim drafts informed by feasibility analysis and prior art.

The 3-Agent Pipeline

Agent 1: Planner --- analyzes the feasibility report and prior art to develop a claim strategy. Identifies the key inventive concepts, determines claim scope, and plans the claim hierarchy (independent claims, dependent claims, method claims, system claims).

Agent 2: Writer --- drafts the actual patent claims following the Planner's strategy. Produces properly formatted independent and dependent claims with correct patent claim language and antecedent basis.

Agent 3: Examiner --- reviews each claim individually, checking for common drafting issues: indefinite language, lack of antecedent basis, overly broad or narrow scope, missing limitations, and consistency with the specification. Provides per-claim feedback.

Claim Features

- **Claim tree visualization** --- SVG-based hierarchical view showing claim dependencies (which dependent claims reference which independent claims), with toggle between list and tree views
- **Claim editing** --- click any claim to edit its text inline
- **Claim regeneration** --- regenerate claims with updated context from feasibility analysis and prior art
- **Word export** --- export claims as a properly formatted Word document
- **Patent family lookup** --- view continuity data (parent applications, child applications, continuations, divisionals) for any prior art patent, fetched from USPTO ODP

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

9. Compliance Checking

The compliance checker validates claim drafts against four key areas of patent law, using 4 specialized LangGraph agents.

The 4 Compliance Checks

Check	Statute	What It Validates
Written Description	35 USC 112(a)	Claims are supported by the specification with sufficient detail
Definiteness	35 USC 112(b)	Claim language is clear and definite, proper antecedent basis
Formalities	MPEP 608	Claims follow USPTO formatting rules and conventions
Eligibility	35 USC 101	Claims are directed to patent-eligible subject matter (Alice/Mayo framework)

Compliance Features

- **Traffic-light results** --- each check produces PASS (green), FAIL (red), or WARN (amber) status
- **MPEP citations** --- specific Manual of Patent Examining Procedure references for each finding
- **Actionable fix suggestions** --- each issue includes a concrete suggestion for how to fix it
- **Per-claim detail** --- compliance results are broken down by individual claim
- **Word export** --- export full compliance results as a formatted Word document

---

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

10. Patent Application Generator

The application generator assembles a complete USPTO-formatted patent application from feasibility analysis, prior art, and claim drafts.

The 5-Agent LangGraph Pipeline

Agent	Section	What It Produces
1	Background	Field of invention, description of related art, problem statement
2	Summary	Brief summary of the invention and its advantages
3	Detailed Description	Full technical description with embodiments and examples
4	Abstract	Patent abstract (150 words or fewer per USPTO rules)
5	Figure Descriptions	Brief description of drawings and figure references

Additional Generated Sections

- **Information Disclosure Statement (IDS)** --- auto-generated table from prior art search results
- **Cross-References** --- user-editable section for related application references

Application Features

- **USPTO-compliant Word export** --- follows 37 CFR 1.52: US Letter, Times New Roman 12pt, 1.5 line spacing, [NNNN] bold paragraph numbering, page numbers, claims and abstract on separate pages
- **DOCX watermark** --- every page includes diagonal "DRAFT --- NOT LEGAL ADVICE --- MUST BE REVIEWED BY PATENT ATTORNEY" watermark
- **Markdown export** --- full application as formatted Markdown with paragraph numbers and IDS table
- **Section navigation** --- 9-section sidebar for easy navigation within the application
- **Inline editing** --- edit any section directly in the browser
- **UPL disclaimer modal** --- shown before application generation to remind users this is a draft, not a filed application

---

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

11. Docker Deployment

PatentForge provides a complete Docker Compose configuration for production-like deployments.

Running with Docker

```
# Set a secure internal service secret
export INTERNAL_SERVICE_SECRET=$(openssl rand -hex 32)

# Optionally set your API key (or configure later in Settings)
export ANTHROPIC_API_KEY=your-key-here

# Build and start all services
docker compose up --build
```

Open http://localhost:8080. Uses PostgreSQL instead of SQLite.

Docker Topology (7 containers)

Container	Service	Port	Notes
frontend	Nginx serving React build	8080	Depends on backend health
backend	NestJS API server	3000	Depends on PostgreSQL
feasibility	Feasibility analysis	3001	Internal only
claim-drafter	Claim drafting	3002	Internal only
application-generator	Application assembly	3003	Internal only
compliance-checker	Compliance validation	3004	Internal only
postgres	PostgreSQL 16	5432	Persistent volume

Docker Features (v0.8.3)

- **Healthchecks** on all 7 services; frontend depends on backend being healthy
- **Log rotation** --- json-file driver with 10 MB max-size, 3-file rotation on all services
- **Data safety** --- no --accept-data-loss flag; schema changes that would drop data fail explicitly
- **Configurable backend port** --- reads PORT from environment (default: 3000)
- **Source maps** at runtime for readable stack traces
- **Startup validation** --- backend validates DATABASE_URL on boot and requires INTERNAL_SERVICE_SECRET in production

---

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

12. Configuration

All settings are configurable via the Settings page in the UI.

Setting	Default	Description
Anthropic API Key	---	Required. Your Claude API key. Encrypted at rest (AES-256-GCM).
USPTO API Key	---	Optional. Free key from data.uspto.gov for enhanced patent search and claims viewing. Encrypted at rest.
Default Model	---	Required. Choose Sonnet, Opus, or Haiku before running analysis. No silent fallback.
Research Model	---	Optional separate model for Stage 2 (e.g., Haiku for cost savings).
Max Tokens	32,000	Maximum tokens per stage response.
Inter-Stage Delay	5 seconds	Pause between stages for rate limit protection.
Cost Cap (USD)	5.00	Enforced server-side. Blocks new runs when cumulative project cost exceeds this amount. Mid-pipeline cancellation when a stage pushes cost over the cap. Set to 0 to disable.
Export Path	Desktop	Folder for saved reports. Must be within your home directory.
Auto Export	true	When enabled, reports are automatically exported on pipeline completion.

Authentication (Optional)

For network deployments, set the PATENTFORGE_TOKEN environment variable to require Bearer token auth on all API requests. When not set, auth is disabled (single-user local mode).

Internal Service Security

The feasibility, claim-drafter, application-generator, and compliance-checker services are internal-only --- the frontend communicates through the NestJS backend, which proxies requests. Set INTERNAL_SERVICE_SECRET to require a shared secret header on all internal service calls.

Docker deployments: The default secret is a known public value. Set a custom secret before deploying:

```
export INTERNAL_SERVICE_SECRET=$(openssl rand -hex 32)
docker compose up --build
```

API Key Rate Limiting (v0.8.3)

The validate-api-key endpoint is limited to 5 requests per 60 seconds via @nestjs/throttler, preventing brute-force API key testing.

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

13. Legal Guardrails

PatentForge includes multiple layers of legal guardrails to ensure users understand the tool's limitations.

Built-In Guardrails

- **First-run clickwrap** --- on first launch, users must acknowledge that PatentForge provides research, not legal advice, before using the app
- **API key disclaimer** --- the Settings page notes that users are connecting to their own Anthropic account and should review the provider's data policies
- **Export watermarks** --- every generated report (HTML, Word, and on-screen) carries a persistent disclaimer stating the output is AI-generated research, not a legal opinion
- **Application UPL disclaimer** --- shown before application generation to remind users the output is a draft for discussion with a patent attorney
- **DOCX watermark** --- every page of exported Word documents includes a diagonal "DRAFT --- NOT LEGAL ADVICE --- MUST BE REVIEWED BY PATENT ATTORNEY" watermark

Security Features

- **API key encryption** --- all API keys (Anthropic, USPTO) encrypted at rest with AES-256-GCM
- **Encryption self-test** --- backend logs ERROR on startup if encrypt/decrypt round-trip fails
- **Helmet security headers** --- X-Content-Type-Options, X-Frame-Options, HSTS, and other headers on all backend responses
- **DOMPurify sanitization** --- server-side HTML sanitization on all marked() outputs prevents XSS in exported reports and DOCX documents
- **Internal service isolation** --- Python services are not directly accessible from the frontend; all requests go through the NestJS backend proxy
- **CORS configuration** --- all services read ALLOWED_ORIGINS env var

---

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

14. What's New in v0.8.3

Released: April 5, 2026

Added

- **Settings: auto-export opt-in** --- new autoExport boolean (defaults to true) with toggle in Settings UI
- **Docker: healthchecks** --- all 7 services now have healthchecks; frontend depends on backend being healthy
- **Docker: log rotation** --- json-file driver with 10 MB max-size, 3-file rotation on all services
- **CI: Node 24 migration** --- FORCE_JAVASCRIPT_ACTIONS_TO_NODE24 env var set preemptively
- **Coverage thresholds bumped** --- backend 44/38/32/43, frontend 38/38/32/38
- **Code quality: shared utilities** --- extracted slugify() and disclaimer text to shared modules

Changed

- **Code quality: lint cleanup** --- all `any` types replaced with proper interfaces (ClaimDrafterResponse, AppGeneratorResponse); unused variables fixed; application.service.ts cost cap now includes claimDraft costs

Fixed

- **CI: package-lock.json version mismatch** --- lockfile version was 0.8.1 while package.json was 0.8.2, causing npm ci to fail in CI
- **decrypt() silent failure** --- decrypt() now throws DecryptionError instead of silently returning ciphertext when decryption fails; settings service shows a clear warning banner in the UI
- **Report auth** --- ReportViewer iframe and HTML download now use fetch + srcdoc/blob instead of direct URLs, working with PATENTFORGE_TOKEN auth
- **Orphaned nginx block** --- removed /feasibility/ proxy block from nginx.conf
- **Stale projectLoadedRef** --- view init ref resets when project ID changes
- **Concurrent run guard** --- isRunningRef prevents overlapping SSE streams from rapid Resume/Cancel/Resume

Security

- **API key rate limiting** --- @nestjs/throttler added; validate-api-key endpoint limited to 5 requests per 60 seconds

---

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

15. What's New in v0.8.2

Released: April 5, 2026

Added

- **Helmet security headers** --- backend now sets security headers (X-Content-Type-Options, X-Frame-Options, HSTS, etc.) via Helmet middleware
- **DOMPurify on all backend HTML generation** --- all marked() call sites sanitized through isomorphic-dompurify, closing the XSS surface on server-rendered HTML
- **ClaimDraft cost tracking** --- claim drafting pipeline costs now saved to database and included in cumulative project cost calculations and cost cap enforcement
- **Application-generator CI job** --- 64 tests now run in GitHub Actions CI alongside other Python services
- **markdownToHtml regression tests** --- 11 unit tests proving the frontend markdown utility escapes script tags, inline event handlers, and raw HTML
- **Claim parser blocklist filter** --- replaced narrow allowlist with a blocklist that rejects known non-claim text while allowing unusual but valid patent claim formats

Changed

- **ProjectDetail.tsx decomposition (phase 2)** --- extracted 6 more units from the coordinator: useViewInit hook, useReportContent hook, ContentPanel, RunningView, ReportView, and StageOutputViewer components; coordinator reduced to 385 lines
- **Prisma strategy consolidation** --- eliminated dual-schema drift by deleting the outdated schema.postgres.prisma and stale migration files; Docker now derives PostgreSQL schema automatically from the canonical schema.prisma
- **npm ci everywhere** --- frontend CI jobs and Dockerfiles switched from npm install to npm ci for deterministic installs
- **Backend build cleanup** --- npm run build now runs rimraf dist && nest build, preventing stale compiled output

Fixed

- **Backend cost-cap tests** --- 8 previously masked cost-cap tests now running and updated
- **Non-claim content in claims** --- items like "Consider adding claims..." now filtered out by blocklist-based validation

Security

- **Helmet** --- security headers on all backend responses
- **DOMPurify** --- server-side HTML sanitization on all marked() outputs

---

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

16. Roadmap

Version	Status	Description
v0.1	Complete	6-stage AI analysis pipeline with streaming
v0.2	Complete	Prior art search, cost tracking, Word export, resume
v0.3.x	Complete	USPTO enrichment, ODP integration, patent claims, E2E tests
v0.4.x	Complete	AI-assisted claim drafting, claim tree, patent family
v0.5.x	Complete	Compliance checking, public release polish
v0.6.x	Complete	Application generator, Docker hardening
v0.7.x	Complete	Windows installer, Mac/Linux beta, system tray, Node SEA
v0.8.x	Complete	Code quality (ESLint, Prettier, strict TS, coverage), security hardening, CI improvements, 595 tests
v0.9	Planned	TBD
v1.0	Planned	Production release

PatentForge is a research tool, not a legal service. Output is AI-generated and must be reviewed by a qualified patent attorney.

17. License

Code (backend, frontend, services): MIT License

Prompt content (services/feasibility/src/prompts/): CC BY-SA 4.0

See LEGAL_NOTICE.md in the repository for important information about this tool's limitations and your responsibilities.

PatentForge v0.8.3 | April 2026 | Scott Converse

https://github.com/scottconverse/patentforge

PatentForge is a research tool, not a legal service. The author is not a lawyer. This tool does not provide legal advice. Output is AI-generated research for personal educational use and must be reviewed by a qualified patent attorney before any filing or legal decision.